

Distributed Tracing in Kubernetes



Piotr Gaczkowski

IT CONSULTANT

@doomhammerng doomhammer.info



Overview



Distributed applications

Passing context for traces

Analyzing traces with Jaeger




Distributed Applications



Even “Simple Web Applications” are distributed




Applications in Kubernetes are distributed by design



Microservices increase the distribution of the components



Distributed Applications



Distributed Tracing


**Allows correlation
of events from
different parts of
the system**

**Particularly fitting
for microservices**

**Helps with
debugging and
optimizing the
code**



Distributed Tracing Model



Source: <https://opentracing.io/docs/overview/>



Trace

A visualization of the life of a request as it moves through a distributed system.



Span



Individual unit of work done in a distributed system



Corresponds to a named and timed operation



Refers to other spans so they can be assembled into one complete Trace




Tags, Logs, and Baggage



Tags are key:value pairs that enable user-defined annotation of spans in order to query, filter, and comprehend trace data



Logs are key:value pairs that are useful for capturing *timed* log messages and other debugging or informational output from the application itself



The **SpanContext** carries data across process boundaries



Baggage Items are key:value pairs that cross process boundaries

OpenTracing

**A standard for distributed tracing
proposed by Jaeger**

Supports many tracers

Supports many languages

Supports plugins for existing solutions

Vendor-neutral API and instrumentation



OpenTracing: Supported Tracers

CNCF Jaeger

LightStep

Instana

Apache SkyWalking

inspectIT

stagemonitor

Datadog

Wavefront by VMware

Elastic APM



OpenTracing: Supported Languages

Go

C++

C#

Java

Javascript

Objective-C

PHP

Python

Ruby



Distributed Tracing Tools

Jaeger

OpenZipkin

Apache
SkyWalking



Jaeger



Cloud Native Computing Foundation
graduated project

Pluggable backend (Elasticsearch,
Cassandra, Kafka)

Native compatibility with OpenTracing

OpenZipkin



Longer in the market

Uses proprietary API or OpenTracing

Hosted by the Apache Foundation



Apache SkyWalking



Hosted by the Apache Foundation

Native compatibility with OpenTracing

Analizes traces and metrics

Instrumenting an Application



Check if the dependencies are already instrumented



Start with areas of value




Cover end-to-end instrumentation



Look for the reuse potential



Passing Context in Process



`start_span` from OpenTracing creates a new span

The `child_of` parameter makes it a child span

This requires passing a span object between function call


In Python, `start_active_span` automatically attaches the span to the parent

Passing Context in Process

```
def main(r=request):  
    span_ctx = tracer.extract(Format.HTTP_HEADERS, request.headers)  
  
    with tracer.start_active_span('processRun', finish_on_close=True,  
child_of=span_ctx) as scope:  
  
        request_id = str(uuid.uuid4())  
  
        app.logger.debug("Received request %s to /", request_id)  
  
        with tracer.start_active_span('validatePayload',  
finish_on_close=True) as scope:  
  
            content = r.json
```



Passing Context for RPC Calls



OpenTracing allows injecting span context into an RPC call

Different RPC protocols are supported including HTTP (REST) and gRPC

Injected context can be extracted on the remote end



Passing Context for RPC Calls

```
with tracer.start_active_span('call downstream', finish_on_close=True) as scope:
```

```
    headers = {}
```

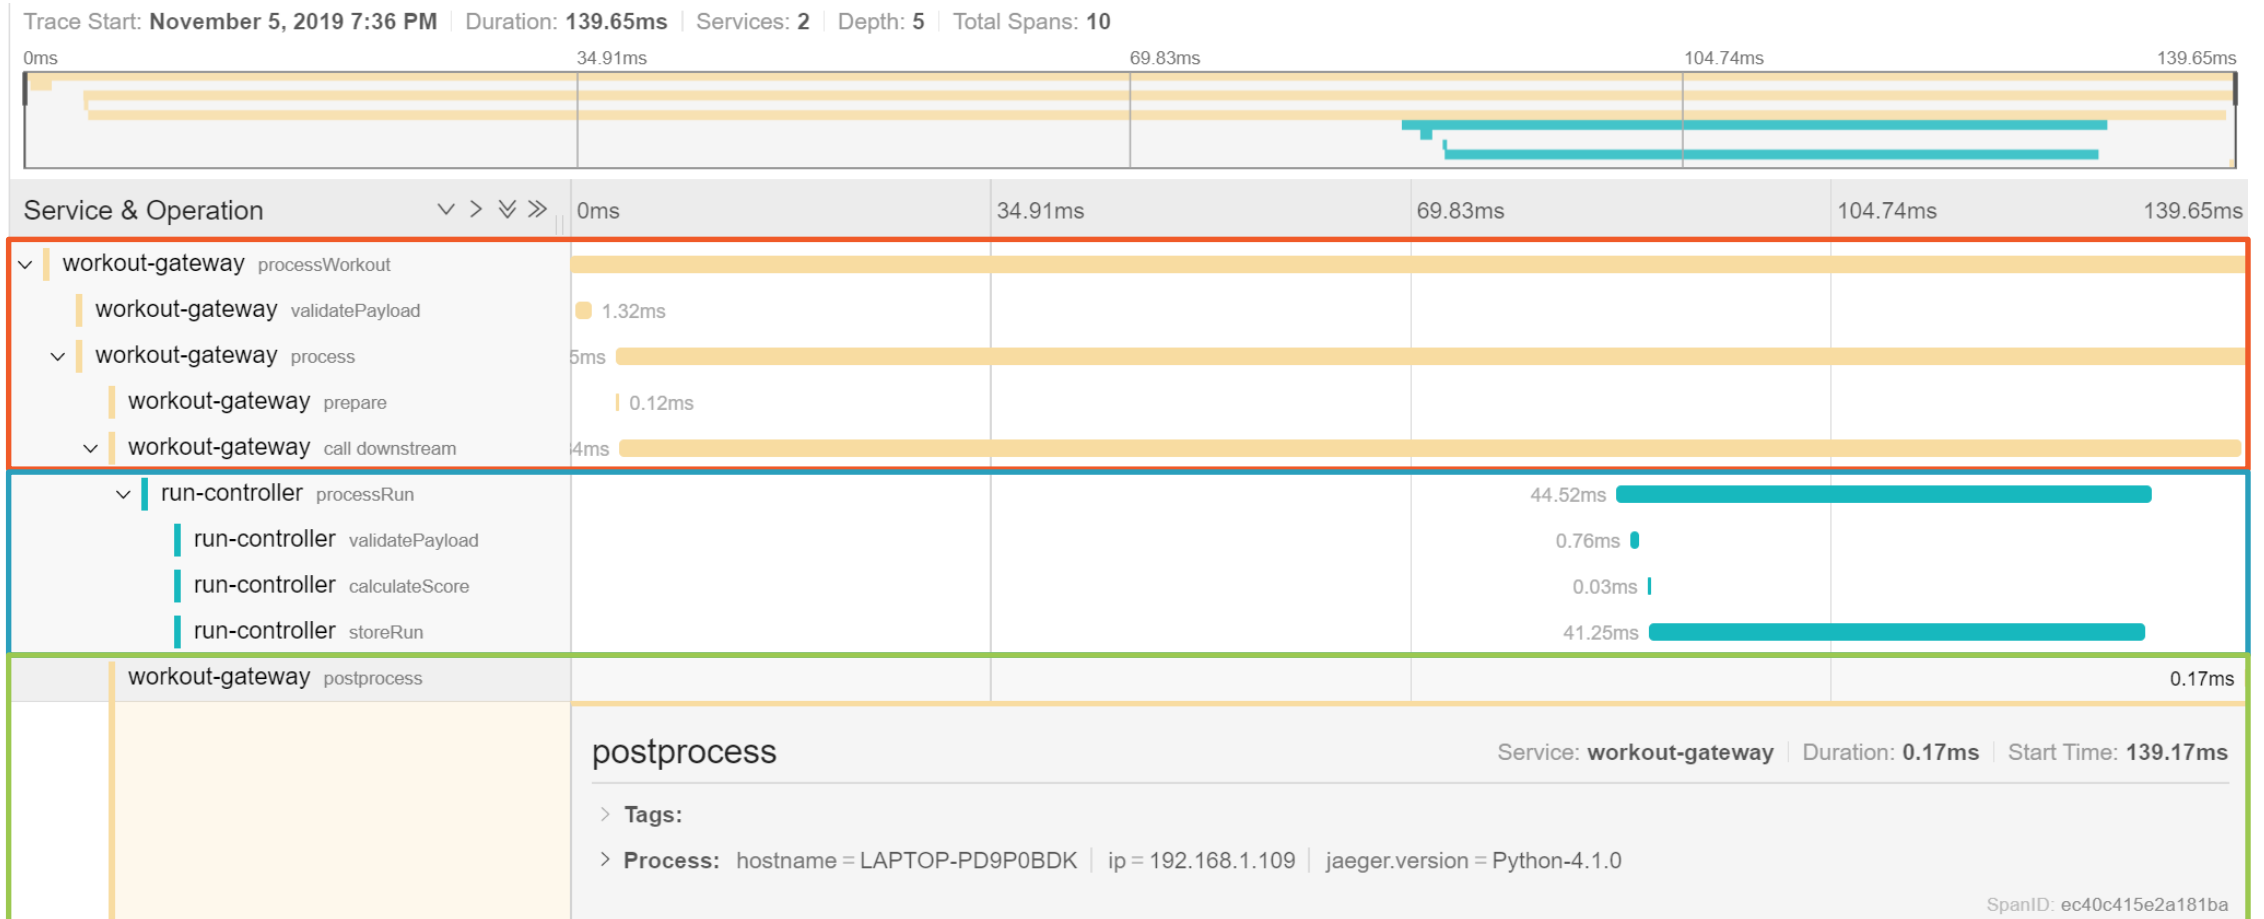
```
    tracer.inject(scope.span, Format.HTTP_HEADERS, headers)
```

```
    payload = {**workout.as_dict(), "workout_id": workout_id}
```

```
    r = requests.post(RUN_CONTROLLER_URL, json=payload, headers=headers)
```



Analyzing Traces with Jaeger



Demo



Deploy Jaeger

Collect traces

Use Jaeger UI to identify performance problems



Summary



Kubernetes applications are distributed by default

Tracing makes it possible to understand the workflow of a distributed system

Precise timing is useful to investigate performance issues

